

CS 440 Assignment 3

Sahil Sharma(226009226)/
Rahul Sankaralingam (212002965)

1.

1. Finite State Space and Absorbing State:

- The game is played on a finite 13×13 , resulting in a total of 169 possible squares.
- The bull and robot occupy separate squares and cannot overlap or occupy obstacle positions, which further restricts the number of possible configurations.
- The target square at (7,7) serves as an *absorbing state*. Once the bull reaches this square, the game ends permanently, as there is no further movement. This structure creates a finite-state Markov process with an absorbing state.

2. Robot's Strategy When Outside the 5×5 Square Around the Bull:

- Let $d_{R,B}$ represent the Manhattan distance between the robot (R) and the bull (B).
- When $d_{R,B} > 2$, the bull moves randomly among its available directions.
- The robot can strategically reduce $d_{R,B}$ by 1 with each move, eventually reaching a position within the 5×5 square around the bull with probability 1. This movement is guaranteed due to the finite state space and the robot's control over its own movement.

3. Robot's Strategy When Inside the 5×5 Square Around the Bull:

- Once the robot is within the 5×5 square surrounding the bull ($d_{R,B} \leq 2$), the bull's movement is biased toward decreasing the distance to the robot.
- By positioning itself strategically, the robot can influence the bull's path toward the target square. Each move the robot makes has a non-zero probability of leading the bull closer to the target, even though the bull moves probabilistically.

4. Robot's Movement Advantage:

- The robot has 8 possible movement directions (up, down, left, right, and diagonals), while the bull is restricted to 4 (up, down, left, right).
- The robot moves first each round, giving it an advantage in controlling the bull's movements.
- By maintaining its position within the 5×5 square around the bull, the robot can systematically lead the bull closer to the target.

Probability-Based Conclusion

Let T represent the time (or number of rounds) it takes for the bull to reach the target square $(7,7)$. We define $T=\infty$ if the bull never reaches the target. Our goal is to show that:

$$\Pr(T=\infty)=0$$

This is supported by the following facts:

- **Finite-State Process:** The game operates within a finite state space with an absorbing state at $(7,7)$.
- **Non-Zero Transition Probability:** The robot's strategic movement ensures that each transition has a non-zero probability of moving the bull closer to the target.
- **Eventual Reachability of the Absorbing State:** Given these conditions, the probability of the bull never reaching the target over an infinite number of moves is zero.

Therefore, the game must end with probability 1. This result relies on the finite-state structure of the game, the probabilistic influence of the robot over the bull's movement, and the presence of an absorbing state at the target.

2. To determine the total number of possible game states, we need to account for the bull's and robot's positions on a 13×13 board while considering constraints such as obstacles, separate positions, and the target.

Steps for Calculating the State Space

1. **Board Configuration and Constraints:**
 - **Board Dimensions:** The board has $13 \times 13 = 169$ squares.
 - **Target Position:** The target square is located at (7,7), where the bull must reach to end the game.
 - **Obstacle Positions:** There are 7 obstacle squares on the board, which are: $\{(6,6), (6,7), (6,8), (7,6), (8,6), (8,7), (8,8)\}$
 - **Initial Positions:**
 - The bull starts at (1,13).
 - The robot starts at (13,1).
 - **Position Constraints:** The bull and robot cannot occupy the same square or any obstacle position at the same time.
2. **Calculating Available Positions:**
 - **Total Squares on the Board:** 169
 - **Unavailable Squares:**
 - 7 obstacle squares
 - 1 square occupied by the other piece (bull or robot)
 - **Available Squares for Each Piece:** $169 - 8 = 161$
3. Thus, each piece (bull and robot) has 161 possible positions on the board at any given time, ensuring they do not overlap and do not occupy obstacle positions.
4. **Calculating the Total Number of Game States:**
 - Since the bull and robot must occupy separate squares, the number of configurations is given by the product of their available positions.
 - **Total States:** Total possible states $= 161 \times 161 = 25,921$

This calculation represents the total possible states for any configuration of the bull and robot on the 13×13 board, considering all constraints.

Conclusion

The total number of valid game states on the 13×13 board, with the given target, obstacles, and position constraints, is 25,921. This finite number of states confirms that the game operates within a limited configuration space.

3.

The goal is to determine the minimum expected number of rounds needed to get the bull to the target. We'll use a formula for $T^*(posB, posC)$, where:

- $posB$ is the bull's position.
 - $posC$ is the robot's position.
 - $T^*(posB, posC)$ is the minimum expected number of rounds to get the bull from $posB$ to the target.
-

Steps to Derive T^*

1. **Base Cases:**

- **Bull at the Target:** If the bull is at (7,7), the game is over, so T^* is 0:

$T^*(posB, posC) = 0$, if $posB = (7,7)$

Bull Trapped but Not at Target: If the bull has no possible moves and isn't at the target, T^* is infinity because it can't reach the target:

$T^*(posB, posC) = \text{infinity}$, if bull is trapped and $posB \neq (7,7)$

Recursive Formula for All Other Cases:

- For other positions, we find $T^*(posB, posC)$ by looking at all possible moves the robot and bull can make.
- The robot can move in 8 directions (up, down, left, right, and diagonals). For each robot move to a new position c , the bull responds by moving to one of its possible positions.
- The recursive formula we use is:

$T^*(posB, posC) = 1 + \min \text{ over } c \text{ in } M_C(posC) [\text{average of } T^*(b, c) \text{ for } b \text{ in } M_B(posB, c)]$

where:

- $M_C(posC)$ is the set of valid moves for the robot from position $posC$.
- $M_B(posB, c)$ is the set of valid moves for the bull from $posB$, given the robot's new position c .

Explanation:

- The formula adds 1 for each round, representing the next move.

- The robot chooses a move c that minimizes T^* .
- We take the average of T^* across all possible bull moves b , assuming the bull responds to the robot's new position c .

Handling Movement Situations:

- The movement pattern changes based on the distance between the robot and bull:

Case 1: Robot is Far from Bull (distance ≥ 2)

In this case, the bull moves randomly, and the robot must work to reduce the distance.

The formula becomes:

$$T^*(\text{posB}, \text{posC}) = 1 + \min_{c \in M_C(\text{posC})} [(1 / |M_B(\text{posB})|) * \text{sum of } T^*(b, c) \text{ for } b \in M_B(\text{posB})]$$

Case 2: Robot is Close to Bull (distance ≤ 2)

Here, the bull is more likely to move toward the robot. $M_B(\text{posB}, c)$ in this case only includes moves that keep or reduce the distance to the robot.

The formula then is:

$$T^*(\text{posB}, \text{posC}) = 1 + \min_{c \in M_C(\text{posC})} [(1 / |M'_B(\text{posB}, c)|) * \text{sum of } T^*(b, c) \text{ for } b \in M'_B(\text{posB}, c)]$$

Robot and Bull Movement Vectors:

- **Robot Moves** (8 directions)

$(dx, dy) \in \{(-1,-1), (-1,0), (-1,1), (0,-1), (0,1), (1,-1), (1,0), (1,1)\}$

Bull Moves (4 directions):

$(dx, dy) \in \{(-1,0), (1,0), (0,-1), (0,1)\}$

Position Constraints:

- Both bull and robot must stay within board boundaries: $1 \leq x, y \leq N$.
- They cannot occupy the same square.

- Neither piece can occupy obstacle positions.

Lower Bound for T^* :

- As a lower bound, $T^*(posB, posC)$ is at least the Manhattan distance from $posB$ to the target $(7,7)$, since that distance is the shortest possible path:

$$T^*(posB, posC) \geq \text{manhattan}(posB, (7,7))$$

Summary

The recursive formula for $T^*(posB, posC)$ considers the robot's best possible moves and the bull's possible responses to ensure the bull reaches the target in the minimum expected time. This formula accounts for all valid moves and uses probabilistic expectations for the bull's movement.

4) Problem 4 Solution: Algorithm for Computing T*

The goal is to compute $T^*(posB, posC)$ which represents the minimal expected number of rounds needed to corral the bull at the target position (7, 7) when the robot starts at position $posC$ and the bull starts at position $posB$.

Algorithm Steps

1. Initialize the Grid and T* Values:

- Define a 2D grid of size $N \times N$ where $N=13$
- For each pair of positions ($posB$, $posC$), initialize $T^*(posB, posC) = \infty$ representing an unsolved state.
- Set $T^*(target, posC) = 0$ for all feasible positions of $posC$, since if the bull is already at the target (7, 7), no rounds are needed for completion.

2. Define Movement Options:

- Define a function $get_robot_moves(posC)$ that returns all 8 possible moves for the robot from position $posC$ (up, down, left, right, and diagonals).
- Define a function $get_bull_moves(posB, posC)$ that returns the 4 possible moves for the bull (up, down, left, right).
 - If the robot is within a 5x5 area centered on the bull, the bull's moves should be adjusted to move closer to the robot when possible.

3. Recursive Update Formula:

- For each position pair ($posB$, $posC$) where $posB \neq target$, calculate $T^*(posB, posC)$ using the recursive formula:

$$T^*(posB, posC) = 1 + \min_{posC' \in get_robot_moves(posC)} \mathbb{E}_{posB' \in get_bull_moves(posB, posC)} [T^*(posB', posC')]$$

4. where:

- $get_robot_moves(posC)$ returns all feasible next positions for the robot, and
- $get_bull_moves(posB, posC)$ returns all feasible next positions for the bull, given the robot's influence within the 5x5 area.
- This formula computes $T^*(posB, posC)$ by finding the minimum number of expected rounds required for each possible robot move $posC'$, assuming the bull takes the move $posB'$ that maximizes the time.

5. Iterate Until Convergence:

- Continuously update $T^*(posB, posC)$ for all grid positions until the values stabilize (i.e., the maximum change across all states is smaller than a convergence threshold ϵ (epsilon)).

- This process ensures that each T^* value represents the minimum expected rounds for the bull to reach $(7, 7)$, given optimal moves by the robot.
6. **Extract the Result:**
- Once the values converge, $T^*(\text{startB}, \text{startC})$ (where startB and startC are the initial positions of the bull and robot, respectively) will provide the minimal expected rounds needed to guide the bull to $(7, 7)$.
-

Explanation of Key Components

- **Initialization:** By setting $T^*(\text{target}, \text{posC})$, we define the absorbing state where the bull reaches the target.
- **Recursive Formula:** The recursive formula balances between minimizing rounds (robot's optimal move) and maximizing rounds (bull's potential response). This simulates an optimal approach for the robot under worst-case bull behavior.
- **Convergence:** Iterating until values stabilize ensures that each T^* value accurately reflects the expected minimum rounds under the defined movement dynamics and constraints.

Final Algorithm in Pseudocode

```
function compute_T_star():
    initialize T_star[(posB, posC)] = infinity for all positions (posB, posC)
    for each posC in grid:
        T_star[(target, posC)] = 0 # Target states initialized to 0

    repeat until convergence:
        max_diff = 0
        for each posB in grid excluding target:
            for each posC in grid:
                min_rounds = infinity
                for each posC' in get_robot_moves(posC):
                    expected_value = 0
                    for each posB' in get_bull_moves(posB, posC):
                        probability = 1 / len(get_bull_moves(posB, posC))
                        expected_value += probability * T_star[(posB', posC')]
                    min_rounds = min(min_rounds, 1 + expected_value)
```



```
# Update T_star and track the maximum difference
max_diff = max(max_diff, abs(T_star[(posB, posC)] - min_rounds))
T_star[(posB, posC)] = min_rounds

return T_star
```

Explanation of Each Step in Pseudocode

- **Initialization:** Sets all positions to *infinity*, except for the absorbing state where $T_star[(target, posC)] = 0$.
- **Outer Loop:** Iterates until the values in *T_star* stabilize to within the convergence threshold.
- **Robot Move Loop:** For each possible move of the robot, computes the expected number of rounds by averaging over all bull moves.
- **Update and Convergence Check:** Tracks the largest change in *T_star* values to determine when to stop iterating.

5) 1. Implementing the Algorithm

- The algorithm developed in Problem 4 was implemented in code to calculate the minimal expected rounds T^* required for the bull to reach the target $(7, 7)$ from any starting positions of the bull and robot.
- We followed the recursive formula and iterative process to update each state's expected rounds, stopping when the values converged based on a convergence threshold $\epsilon=0.01$

2. Running the Algorithm

- Starting Positions: The bull starts at $(1, 12)$ and the robot at $(12, 0)$.
- Convergence Threshold: We used an epsilon value of 0.01 to ensure stability and precision in the expected rounds calculation.
- Result Consistency: The code was run multiple times and consistently returned the same result.

3. Results

- Expected Time: The code outputs an expected time of 18.11 rounds.
- Interpretation:
 - Since the expected time is finite, this confirms that the robot can indeed pen the bull from the initial positions within finite time.
 - The result, 18.11 rounds, represents the minimum expected number of rounds needed for the robot to guide the bull to the target, considering both random and directed movements of the bull based on the robot's influence.

4. Final Answer

- Can the robot pen the bull in finite expected time?
 - Yes, the robot can pen the bull in finite time.
- What is the expected time?
 - The expected time for the bull to reach the target is approximately 18.11 rounds.

6)

1. Robot's Diagonal Advantage:
 - Moving diagonally allows the robot to reach positions that the bull can't access in a single move, giving the robot more options for avoiding direct capture.
 - Diagonal movement also allows the robot to better position itself around the bull's influence zone, potentially guiding the bull toward (7, 7) without being directly in its path.
2. Influence of Obstacles Around the Target:
 - Obstacles near (7, 7) limit entry points and safe positioning for the robot as it approaches the target.
 - To pen the bull, the robot must guide it close to (7, 7)—but in doing so, the robot risks getting cornered by obstacles and reducing its movement options.
3. Bull's Direct Chase:
 - The bull's aggressive movement will always prioritize reducing the Manhattan distance to the robot, which puts constant pressure on the robot.
 - The bull only needs to get within one move of the robot orthogonally to crush it, so the robot needs to stay at least two cells away at all times. This becomes difficult around the target, where obstacles reduce maneuverability.

Final Determination: Can the Robot Win?

Considering the grid constraints and obstacle layout, here's the outcome breakdown:

- Early Game (Starting Positions): The robot has space and can use its diagonal advantage to influence the bull's path from a safe distance.
- Mid-Game (Guiding the Bull Toward the Target): The robot can continue guiding the bull effectively by positioning itself strategically in the 5x5 influence zone without getting too close.
- Endgame (Near the Target):
 - As the bull approaches (7, 7), the obstacles create a bottleneck for the robot. The robot's movement becomes more restricted as it tries to influence the bull's final moves without getting too close.
 - The limited safe zones near (7, 7) mean the robot may not be able to maintain a safe two-cell distance, allowing the bull to close in.

Conclusion

Even with diagonal movement, the robot is likely unable to maintain both influence and safety as it brings the bull close to (7, 7). The constraints around the target favor the bull's aggressive chase, and in optimal play:

- The bull is likely to win, as the robot will run out of safe spaces while trying to guide the

bull to the target.

7) Objective

We want a heuristic that estimates $T^*(\text{posB}, \text{posC})$, which represents the minimum expected rounds to guide the bull to $(7, 7)$ from any starting configuration of the bull and the robot.

Key Characteristics of a Good Heuristic

To be effective, the heuristic should:

1. **Never overestimate** the true T^* (admissibility), ensuring that it's optimistic.
2. **Approximate the remaining rounds needed** with minimal computational overhead.
3. Be **consistent** if possible, meaning it satisfies the triangle inequality.

Proposed Heuristic: Manhattan Distance from Bull to Target

A practical and straightforward heuristic for $T^*(\text{posB}, \text{posC})$ is the **Manhattan distance** from the bull's current position posB to the target $(7, 7)$.

$$h(\text{posB}) = \text{abs}(x_{\text{posB}} - 7) + \text{abs}(y_{\text{posB}} - 7)$$

where x_{posB} and y_{posB} represent the coordinates of the bull's position.

Why This Heuristic is a Good Choice

1. **Admissibility:**
 - The Manhattan distance heuristic represents the minimum number of orthogonal moves required to reach the target $(7, 7)$.
 - Since each move reduces the distance by at most 1, this heuristic provides a lower bound on the expected rounds, making it admissible.

2. Consistency:

- The Manhattan distance satisfies the triangle inequality. For any transition from $posB$ to a neighboring position $posB'$, the heuristic estimate $h(posB)$ will not exceed the cost of moving to $posB'$ plus the heuristic $h(posB')$
- This consistency is important for search algorithms like A* because it ensures that the heuristic guides the search effectively without overestimating costs.

3. Practical Usefulness:

- **Guiding the Robot's Strategy:** Using $h(posB)$ helps the robot plan its movements in a way that minimizes the expected rounds. For example, the robot can stay within the bull's influence zone to gradually reduce the Manhattan distance to the target.
- **Computational Efficiency:** Calculating the Manhattan distance is computationally simple, requiring only addition and subtraction. This makes it feasible to apply this heuristic repeatedly during a search or optimization process without significant performance cost.

4. Influencing the Bull's Movement:

- As the robot's objective is to guide the bull, $h(posB)$ can also inform when the robot should adjust its position to lead the bull closer to the target.

Final Answer

- **Heuristic:** The Manhattan distance from the bull's position $posB$ to the target $(7, 7)$:
 $h(posB) = \text{abs}(x_posB - 7) + \text{abs}(y_posB - 7)$
Usefulness: This heuristic provides an admissible and consistent estimate for T*helping the robot navigate optimally to pen the bull while minimizing computational complexity.